



A recommender system based on customer interaction with an e-commerce website

Richard Hruby (17-619-172)
Johan Faxner (21-603-204)
Tim Matheis (21-603-907)
Giovanni Magagnin (17-300-914)

Thursday, 12th of May 2022. Big Data Analytics (8,727,1.00)

Motivation


The number of digital buyers worldwide went from 1.32b in 2014 to 2.14b in 2021 (Statista)

The availability of information increases the importance of systematic data collection, data processing, data analysis and implementation in various business and value chains


Several studies (inter alia: Accenture, 2018; PwC, 2019; Deloitte, 2019) indicate consumers' wish and demand for personalization

➔ How well can we predict the purchases of an online shopper based on their previous shopping behavior as well as the purchasing behavior of other customers?

**E-COMMERCE AS SHARE OF TOTAL
RETAIL SALES WORLDWIDE 2015-2025**



**LEADING SOURCES OF INSPIRATION FOR
ONLINE SHOPPERS WORLDWIDE AS OF
APRIL 2021**



Methodology

Recommender system

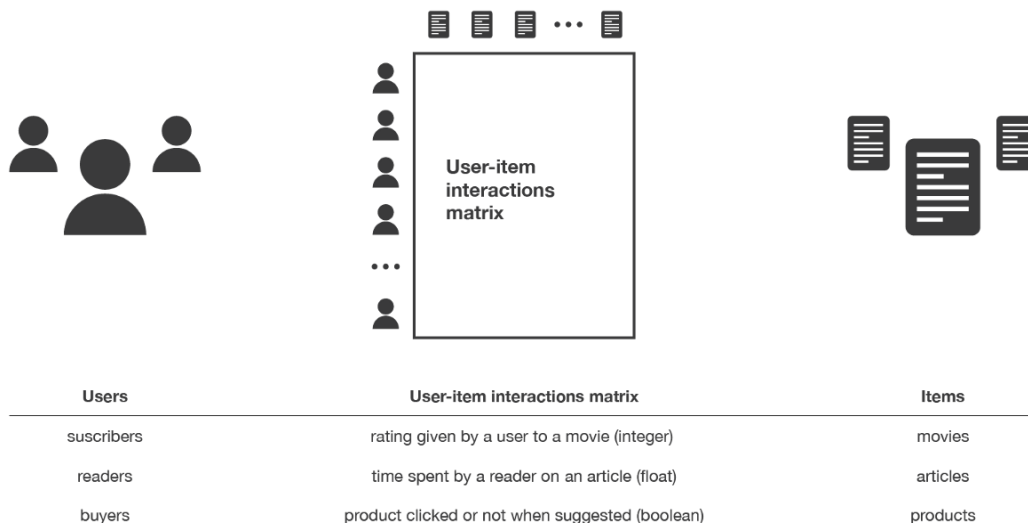
Subclass of Information filtering Systems that seeks to predict the rating or the preference a user might give to an item.

Requirements:

- A set of products and a set of customers who, in some way, have interacted with these products.

User-based collaborative filtering

- Past users' behaviour to generate future preferences
- We discover the relevant features based on patterns
- The features are not “human-based” they are the result of an algorithm.



Alternating least squares (ALS)

- Suitable for larger-scale collaborative filtering problems
- ALS minimizes two loss functions alternatively
- ALS runs its gradient descent in parallel across multiple partitions of the training data

Item

W X Y Z

A

B

C

D

User

Interaction Matrix

=

A

B

C

D

User Matrix

X

W X Y Z










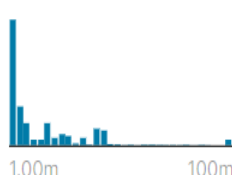
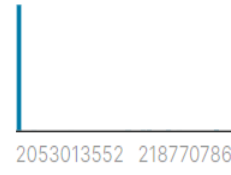
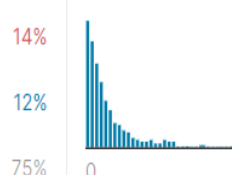
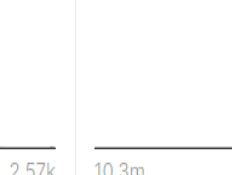
1.5 1.2 1.0 0.8

1.7 0.6 1.1 0.4

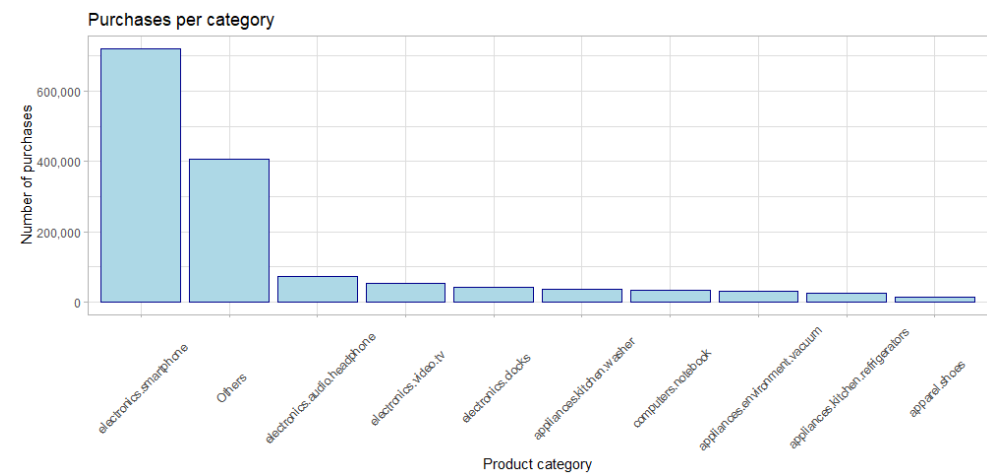
Item Matrix

Data-set

We used a 15GB data-set with behaviour data for 2 months from a large multi-category online store from Kaggle. The data is organized in the following way:

 event_time	 event_type	 product_id	 category_id	 category_code	 brand	 price	 user_id	 user_session
When event is was happened (UTC)	Event type: one of [view, cart, remove_from_cart, purchase]	Product ID	Product category ID	Category meaningful name (if present)	Brand name in lower case (if present)	Product price	Permanent user ID	User session ID
67501979 total values	view 94% cart 4% Other (916939) 1%			[null] 32% electronics.smartp... 24% Other (29228808) 43%	[null] 14% samsung 12% Other (50394499) 75%			13776051 unique values
		1.00m100mb	20530135522187707861b			02.57k	10.3m580m	

10 categories make up the most purchases in the set. Smartphones and “others” account for most sales.



Most active customers that buy at least one product, also only buy one product.

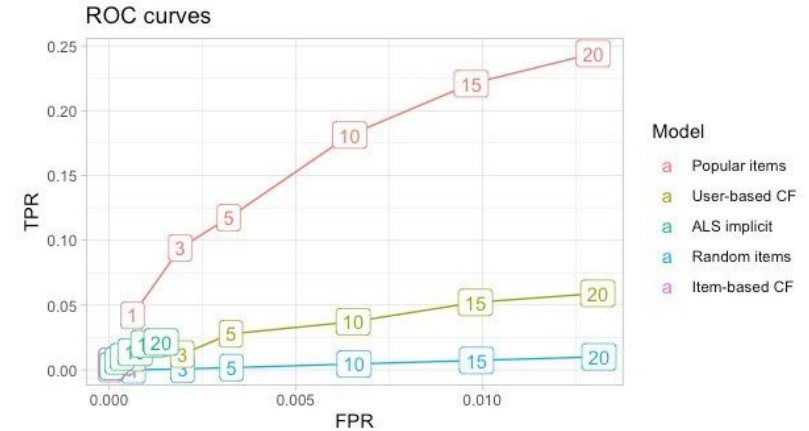


Implications of the big data set – Issues we faced while working with 15GB

Started with a subset of our data (200k rows):

- We used recommenderlab, R package that provides an infrastructure to test and develop recommender algorithms

➔ With the reduced data-set of 200k rows, it worked fine



We tried the same approach with the entire data, but found 2 bottlenecks:

- **Bottleneck 1:** User-item matrix grows exponentially

```
> # spread df, set non-existent values to 0 to use binary encoding
> matrix_data <- cosine_user_product_df %>%
+   select(user_id, product_id, number_purchases) %>%
+   spread(product_id, number_purchases, fill=0)
Error: cannot allocate vector of size 353.8 Gb
```

number of rows	memory usage
200k	35 MB
400k	124 MB
800k	362 MB

- **Bottleneck 2:** Larger user-item matrix implies slower predictions

```
> system.time(recom_topNList2 <- recommenderlab::predict(object = r2,
+   newdata = tail(ratings_matrix, n_pred),
+   type = "topNList",
+   n = n_rec)
+ )
```

number of rows	time elapsed
200k	13.3 sec.
400k	41.5 sec.
800k	93.5 sec.

Big data techniques – How we solved the issues we were facing



```
# purchases per user of all products
products_purchased_per_customer <-|
"
  SELECT user_id, product_id, 1 AS purchased
  FROM ecom
  WHERE event_type IN ('purchase')
  ORDER BY user_id ASC;
"

# use query to extract needed data
purchase_df <- dbGetQuery(ecom_db, products_purchased_per_customer)
purchase_df
```



Shift memory load from RAM to SDD



Split CSV file into tables in SQL



Enable to query important information



```
# configure spark connection
config <- spark_config()
config$spark.executor.memory <- "8G"
config$spark.executor.cores <- 4
config$spark.executor.instances <- 3
config$spark.dynamicAllocation.enabled <- "false"

# initiate connection
sc <- spark_connect(master = "local", config=config,
  version = "3.1")
```



Distribute computationally expensive processes



Migrate pipeline to SparklyR



Enable scaling to the entire dataset



```
# Parallelize revenue calculation
ncores <- parallel::detectCores()
ctemp <- makeCluster(ncores)

system.time(
  output <- foreach(i = c(1, 3, 5, 10, 15, 20, 25), .combine = rbind)%dopar%{
    revenues_n = generate_revenue_df(i, predicted)
    rev_n = list(name = "ALS Spark", n = i, revenue = sum(revenues_n$revenue))
  }
)["elapsed"] #Parallelized calculation is slightly faster

# Unpack results to DF
rev_df = data.frame(output) %>% `rownames<-`(1:7) %>%
  unnest(cols = c(name, n, revenue))
```



Speed up training and evaluation of recommender



Joins and parallelized loops




Enable efficient processing and real-time predictions

Performance of big data methods – Speed and efficiency improvements

Recommenderlab

Inefficient:

- Small rating matrix size (2872 x 1523)
- Slow prediction time: ~15 seconds for 1000 users
- Large memory allocation: used entire RAM on laptop




```
> bench::mark(
+   predict_fun(r2, ratings_matrix, n_pred, n_rec)
+ )
# A tibble: 1 × 13
  expression                min   median `itr/sec` mem_alloc
  <bch:expr>          <bch:tm> <bch:tm>      <dbl> <bch:byt>
1 predict_fun(r2, ratings_matrix, n_pred, n_rec)    15.4s    15.4s    0.0649    7.55GB
# ... with 8 more variables: `gc/sec` <dbl>, n_itr <int>, n_gc <dbl>,
#   total_time <bch:tm>, result <list>, memory <list>, time <list>, gc <list>
```

Sparklyr

Efficient:

- Large rating matrix size (48187 x 9015)
- Same prediction time, but for a matrix 100 times the size.
- Small memory allocation: negligible RAM required



```
> bench::mark(
+   ml_predict(als_fit, spark_test_complete) %>% collect()
+ )
# A tibble: 1 × 13
  expression                min   median `itr/sec` mem_alloc
  <bch:expr>          <bch:tm> <bch:tm>      <dbl> <bch:byt>
1 ml_predict(als_fit, spark_test_complete) %>% collect()    15.1s    15.1s    0.0664    89.6MB
# ... with 7 more variables: n_itr <int>, n_gc <dbl>, total_time <bch:tm>, result <list>,
#   memory <list>, time <list>, gc <list>
```

Results of the prediction – Our assesment

1. Top-N predictions for 2 sample users

	user_id	product_id	purchased	prediction	
	<int>	<int>	<int>	<dbl>	
User 1	1	512503671	1004856	1	0.490
	2	512503671	1004767	0	0.447
	3	512503671	1004833	1	0.385
User 2	4	513997627	1005115	1	1.02
	5	513997627	1004856	0	0.937
	6	513997627	1004767	0	0.859

n = 3

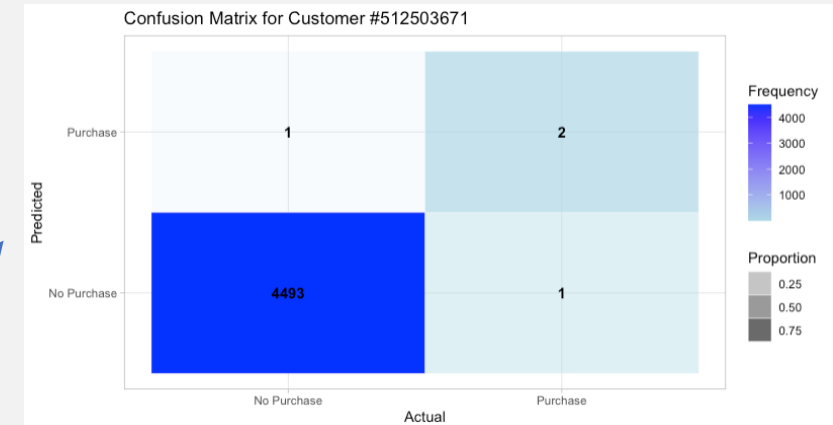
2. Evaluation Metrics per User

	user_id	actual_positives	actual_negatives	true_positives	false_positives	true_negatives	false_negatives
	<int>	<int>	<int>	<int>	<int>	<int>	<int>
User 1	1	512503671	3	4494	2	1	4493
User 2	2	513997627	1	4496	1	2	4494

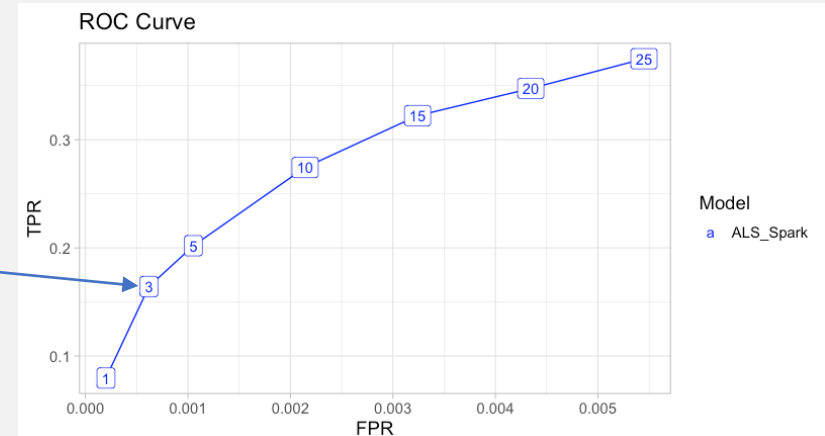
3. Average Evaluation Metrics for N 1:25

	name	n	TPR	FPR	precision	recall	conversion
1000 users	ALS_Spark	3	0.16422619	0.0006199274	0.07100	0.16422619	0.07100

This customer purchased 2 out of 3 recommended products



With 15 recommendations we capture 1/3 of the products purchased



Business and economic implications

Business value

Real increase in revenues:

- With n (recommendations x user) = 5, the business has a net increase of >100k USD
- With n (recommendations x user) = 15, the business has a net increase of >200k USD

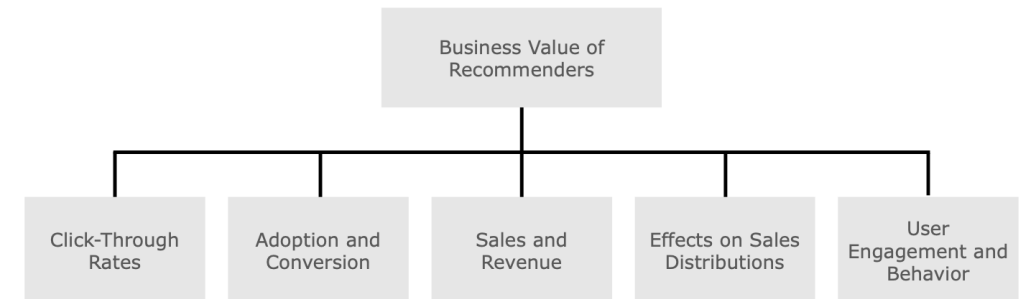
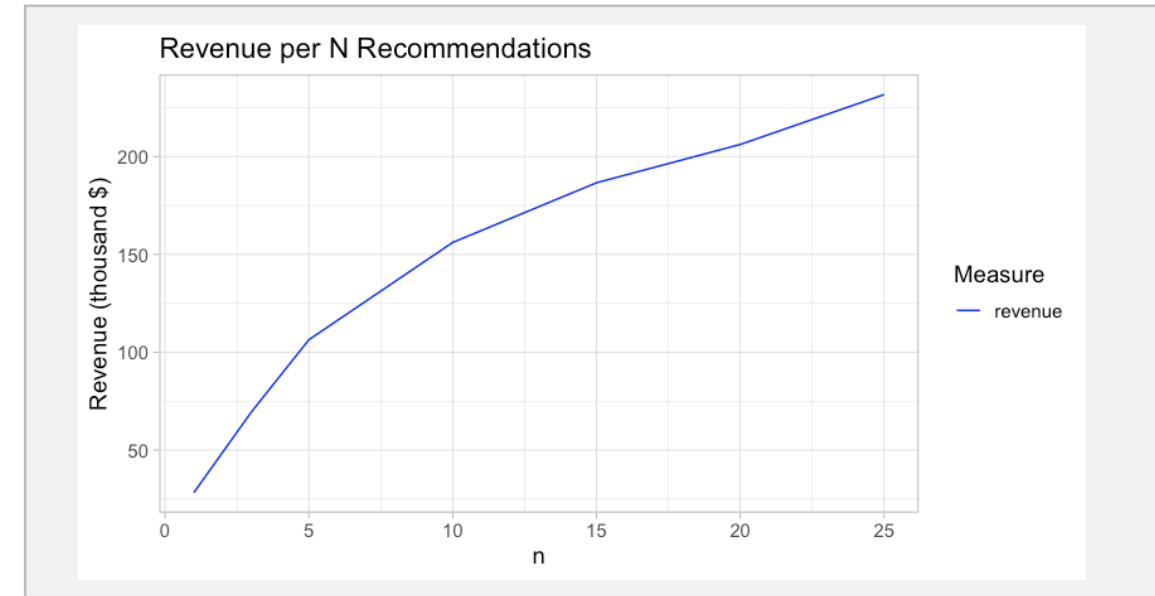
Economic implications

Positive welfare for sellers (Fleder and Hosanagar 2009):

1. Meeting the consumers' individual preferences
2. Intensifying the users' transactions and consumption time
3. Turning browsing consumers into buyers, cross selling, increasing customer loyalty

Positive welfare for buyers:

1. Transaction costs can be reduced, more precisely search costs
2. Solves the information overload problem



Source: Jannach & Jugovac (2019)

Limitations of our approach – What else could be done?



Limitations:

1. **Backward looking bias:** We test whether we would have recommended a product to a customer that actually bought this product. However, the «wrong» predictions could have also led to purchases when the product would have been recommended to the customer.
2. **Methods:** simplistic «market-basket analysis»
3. **Limited data:** only two months of data used for analysis
4. **Product/Customer differentiation:** We treat the products/customers as if they were equal. But it would be useful to include the characteristics of the products.



Future work:

1. **Measure real-life business value:** AB testing to figure out how well recommendations work
2. **Customer behaviour:** Including the other interactions of users for predictions (e.g. view, put in cart, and remove from cart)
3. **Scaling up:** AWS for more robust data pipeline and quicker computation
4. **Hybrid recommender systems:** Incorporate further information about products and users to increase relevance of recommendations

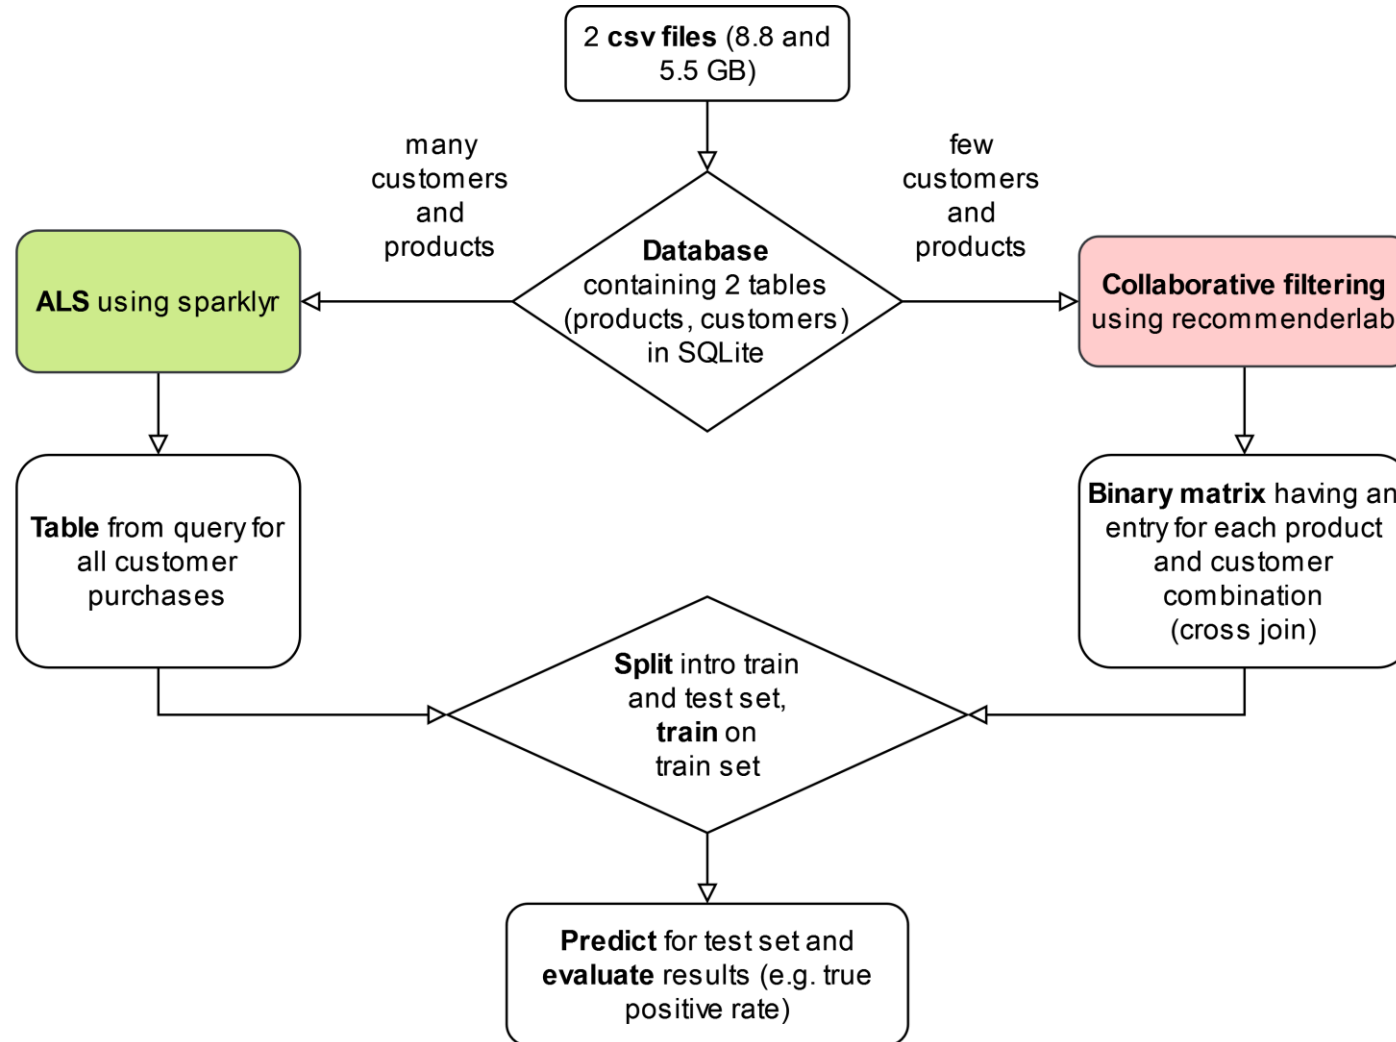


A recommender system based on customer interaction with an e-commerce website

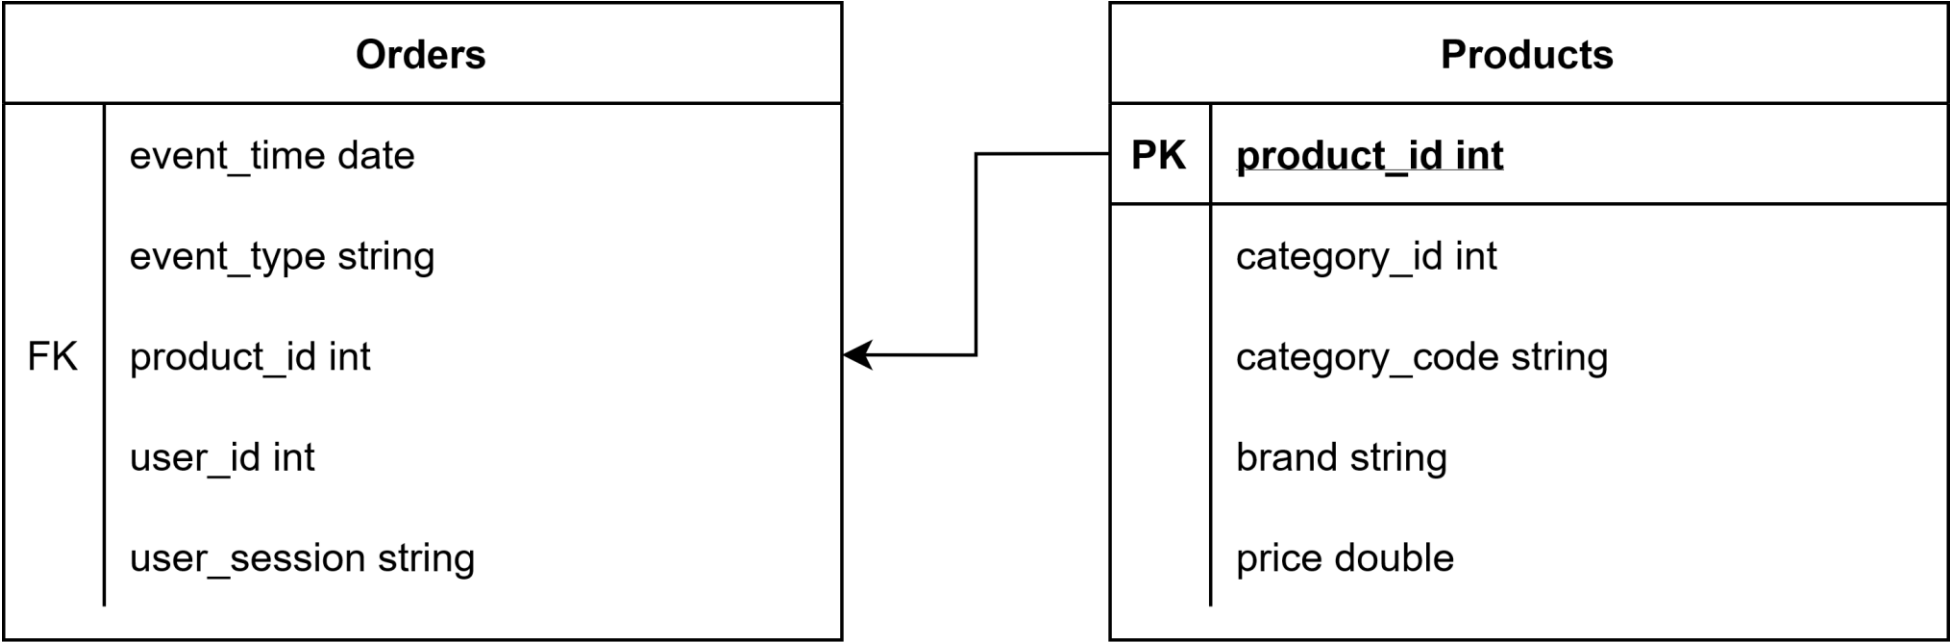
Richard Hruby (17-619-172)
Johan Faxner (21-603-204)
Tim Matheis (21-603-907)
Giovanni Magagnin (17-300-914)

Thursday, 12^o of May 2022. Big Data Analytics (8,727,1.00)


Discussion



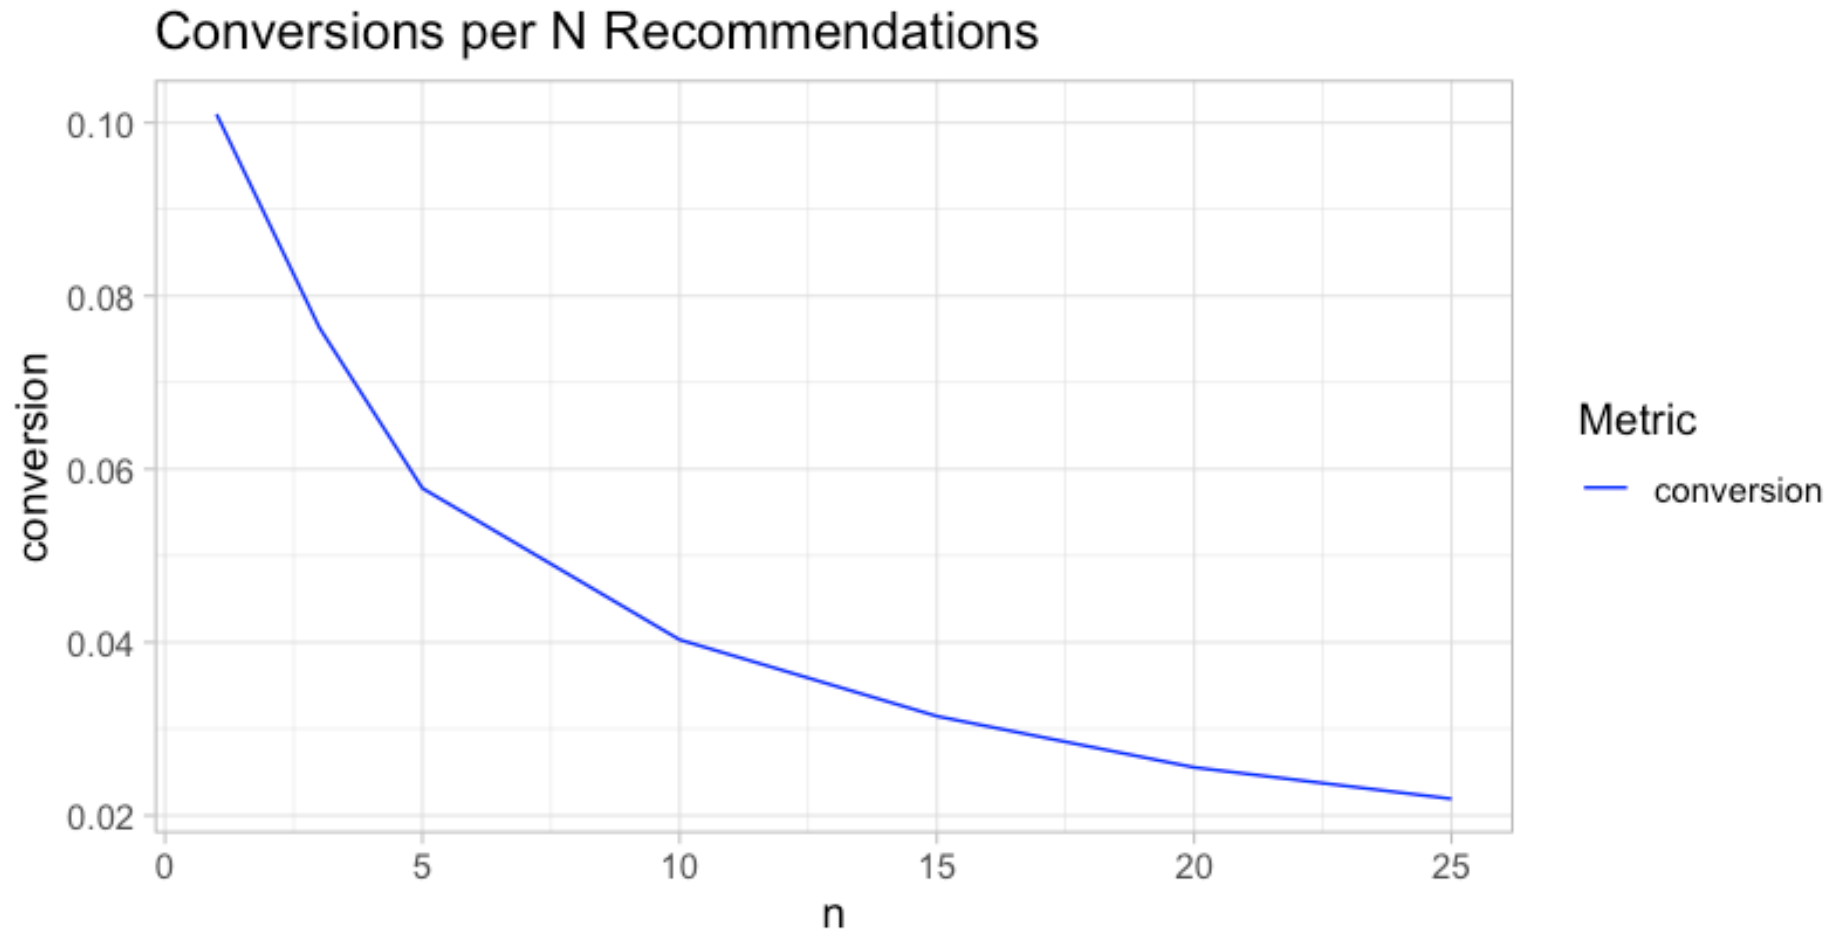
Appendix 1 - SQL



Appendix 2 – Precision Recall Curve



Appendix 3 – Conversions per N Recommendations



Appendix 3 – Top-N Evaluation N=1:25

	name	n	TPR	FPR	precision	recall	conversion	revenue
1	ALS_Spark	1	0.07966667	0.0001999695	0.10100	0.07966667	0.10100	<u>28255</u>
2	ALS_Spark	3	0.16422619	0.0006199274	0.07100	0.16422619	0.07100	<u>69463</u>
3	ALS_Spark	5	0.20195238	0.0010532321	0.05300	0.20195238	0.05300	<u>106393</u>
4	ALS_Spark	10	0.27434405	0.0021429452	0.03660	0.27434405	0.03660	<u>156102</u>
5	ALS_Spark	15	0.32221310	0.0032411125	0.02860	0.32221310	0.02860	<u>186594</u>
6	ALS_Spark	20	0.34742143	0.0043455085	0.02320	0.34742143	0.02320	<u>206173</u>
7	ALS_Spark	25	0.37484762	0.0054499049	0.01996	0.37484762	0.01996	<u>231659</u>

References

- <https://gist.github.com/twolodzko/7becd98ff256ef826b56945de297700d>
- <https://utd-ir.tdl.org/bitstream/handle/10735.1/5485/ETD-5608-7403.99.pdf?sequence=5&isAllowed=y>
- <http://www.learconference2015.com/wp-content/uploads/2014/11/Calvano-slides.pdf>
- <https://www.econstor.eu/bitstream/10419/228752/1/174515275X.pdf>
- <https://towardsdatascience.com/introduction-to-recommender-systems-6c66cf15ada>
- <https://towardsdatascience.com/prototyping-a-recommender-system-step-by-step-part-2-alternating-least-square-als-matrix-4a76c58714a1>